

Little Free Library Analysis

Kaleb Crans

2023-04-17

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.1    v purrr   1.0.1
## v tibble  3.1.8    v dplyr   1.1.0
## v tidyr   1.3.0    v stringr 1.5.0
## v readr   2.1.4    v forcats 1.0.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(readr)
library(sf)

## Linking to GEOS 3.11.0, GDAL 3.5.3, PROJ 9.1.0; sf_use_s2() is TRUE

library(spData)

## To access larger datasets in this package, install the spDataLarge
## package with: `install.packages('spDataLarge',
## repos='https://nowosad.github.io/drat/', type='source')`

rm(list = ls())

if (file.exists("lfl.RData")) {
  load("lfl.RData")
} else {
  libraries = as_tibble(read_csv("libraries.csv"))
  save(libraries, file = "lfl.RData")
}
```

Data Cleaning and Preparation

Let's take a look at the different data types of the columns:

```
lapply(libraries, typeof)

## $id
## [1] "double"
##
## $Name
## [1] "character"
##
## $Street__c
## [1] "character"
##
```

```
## $City__c
## [1] "character"
##
## $State_Province_Region__c
## [1] "character"
##
## $Postal_Zip_Code__c
## [1] "character"
##
## $Country__c
## [1] "character"
##
## $Traveling_Library__c
## [1] "logical"
##
## $Official_Charter_Number__c
## [1] "character"
##
## $First_Map_Date__c
## [1] "double"
##
## $Map_Me__c
## [1] "character"
##
## $Map_Date__c
## [1] "double"
##
## $Duplicate_Charter_Number__c
## [1] "logical"
##
## $Count_of_Primary_Stewards__c
## [1] "double"
##
## $Latitude_MapAnything__c
## [1] "double"
##
## $Longitude_MapAnything__c
## [1] "double"
##
## $Library_Geolocation__Latitude__s
## [1] "double"
##
## $Library_Geolocation__Longitude__s
## [1] "double"
##
## $check_in_count
## [1] "double"
```

Map_Me__c needs to be transformed into a logical variable, as “Taken Down Temporarily” and “Mapped” are the only two categories.

```
libraries <- libraries %>% mutate(Map_Me__c = Map_Me__c == "Mapped")
```

Note that there is only one row with this flag set to false:

```
sum(libraries$Map_Me_c == FALSE)
```

```
## [1] 1
```

A quick look at `check_in_counts` shows that this feature is rarely used when you consider how many times a given library is actually visited:

```
print("Max:")
```

```
## [1] "Max:"
```

```
max(libraries$check_in_count)
```

```
## [1] 47
```

```
print("Summary Stats:")
```

```
## [1] "Summary Stats:"
```

```
summary(libraries$check_in_count)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
##    0.000  0.000   1.000   1.567   2.000  47.000
```

```
check_in_dist <- libraries %>% count(check_in_count)
```

```
annotations <- tibble(x = c(min(check_in_dist$check_in_count), median(check_in_dist$check_in_count), m
```

```
                        y = c(27000, 1000, 1000),
```

```
                        label = c("Min:", "Median:", "Max:"))
```

```
ggplot(data = libraries, aes(x = check_in_count)) +
```

```
  ggtitle("Little Free Library check-in count distribution") +
```

```
  xlab("Number of check-ins") +
```

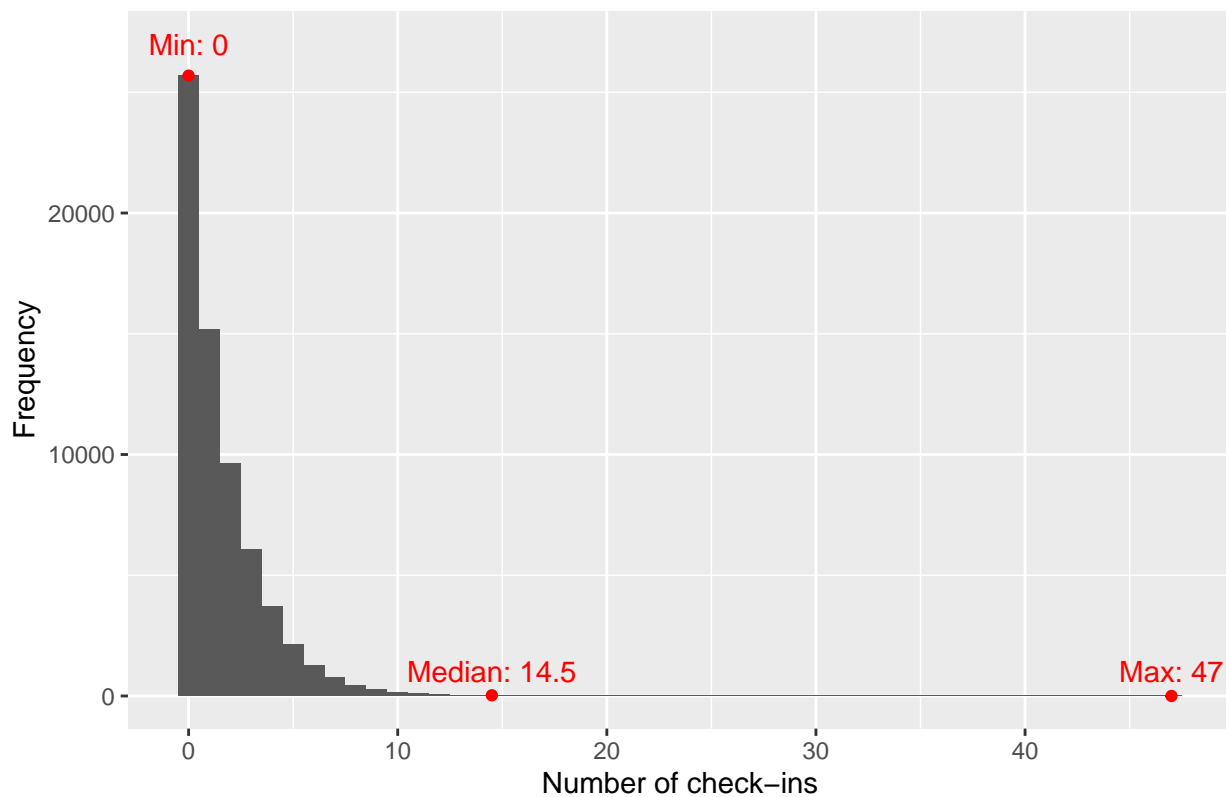
```
  ylab("Frequency") +
```

```
  geom_histogram(binwidth = 1) +
```

```
  geom_point(data = tibble(x = annotations$x, y = c(25691, mean(c(41, 17)), 1)), aes(x = x, y = y), col
```

```
  geom_text(data = annotations, aes(x = x, y = y, label = paste(label, x)), size = 4, color = "red")
```

Little Free Library check-in count distribution



```
#ggsave("checkins.png")
```

```
# Clean up alternative names for the same country:
```

```
libraries <- libraries %>% mutate(Country__c = replace(Country__c, Country__c %in% c("USA", "US", "U.S.",
```

```
# Ranking of countries by number of little free libraries:
```

```
libraries %>% count(Country__c) %>% arrange(desc(n)) %>% head(10)
```

```
##      Country__c      n
## 1   United States 59816
## 2      Canada    2871
## 3      <NA>      633
## 4      Italy     367
## 5  United Kingdom  278
## 6    Australia   229
## 7     France    202
## 8  Philippines   123
## 9  Netherlands   108
## 10   Belgium     90
```

The U.S. by and large has the greatest amount of little free libraries (with alternate spellings outpacing many countries even). Canada is the only country with a somewhat comparable amount, specifically if you adjust for population size.

```
# US population
```

```
us_pop <- 331900000
```

```
# Canada population
```

```
can_pop <- 38250000
```

```
us_count <- libraries %>% filter(Country__c == "United States") %>% nrow
can_count <- libraries %>% filter(Country__c == "Canada") %>% nrow
```

So the per capita number of little free libraries in the US is:

```
format(us_count/us_pop, scientific = FALSE)
```

```
## [1] "0.000180223"
```

And in Canada is:

```
format(can_count/can_pop, scientific = FALSE)
```

```
## [1] "0.00007505882"
```

So we can conclude that Little Free Libraries are a predominately American phenomenon. For the purposes of further analysis, let's exclude all data points not in the US. We will also include rows that have a NA country value (but still a state code) as on inspection many of them are actually located in the US. We can impute the correct country code based off the coordinates later.

```
libraries <- libraries %>% filter(Country__c == "United States" | (is.na(Country__c) & !is.na(State_Prov
```

Analysis by state

How about the distribution by state?

```
length(unique(libraries$State_Province_Region__c))
```

```
## [1] 174
```

But there's only 50 states! So we need to do some data cleaning first.

```
libraries %>% count(State_Province_Region__c) %>% arrange(desc(n)) %>% tail(20)
```

```
##      State_Province_Region__c n
## 155                USA 1
## 156                  Ut 1
## 157             Wyoming 1
## 158                   X 1
## 159                  co 1
## 160                  ct 1
## 161             kentucky 1
## 162                   ks 1
## 163                   ma 1
## 164                   md 1
## 165                   mi 1
## 166                   mo 1
## 167                   ny 1
## 168                   ohio 1
## 169                   ok 1
## 170             omaha 1
## 171                   pa 1
## 172                   tn 1
## 173                   va 1
## 174                   wv 1
```

There's a bunch of different spelling variations. Let's instead take the actual coordinates and then find the states ourselves (assuming the coordinates are correct). One point of interest in the dataset to note is that

there are two sets of coordinates for each row: Latitude_MapAnything__c and Longitude_MapAnything__c vs Library_Geolocation__Latitude__s and Library_Geolocation__Longitude__s.

We can make a dataframe with the differences as separate columns, and print out the mean difference in latitude and longitude respectively:

```
differences <- libraries %>% mutate(dif_lat = (abs(Latitude_MapAnything__c) - abs(Library_Geolocation__Latitude__s)),
  dif_long = (abs(Longitude_MapAnything__c) - abs(Library_Geolocation__Longitude__s)))
```

```
## [1] -1.502292 -3.769037
```

Unfortunately due to the curvature of the earth these values as-is don't reflect physical distances

To explore further, let's take one example with a latitude difference of 16.6 and then plug the coordinates into Google maps. We get two different locations, one in Lake Park Iowa and the other in Lake Park Florida:

```
libraries %>% filter(Latitude_MapAnything__c == 26.79489)
```

```
##      id      Name  Street__c  City__c State_Province_Region__c
## 1 41210 LIB-000004180 307 4th St. Lake Park                      FL
##   Postal_Zip_Code__c  Country__c  Traveling_Library__c
## 1          33403 United States                      FALSE
##   Official_Charter_Number__c First_Map_Date__c Map_Me__c Map_Date__c
## 1          11455      2014-10-03      TRUE  2019-01-22
##   Duplicate_Charter_Number__c Count_of_Primary_Stewards__c
## 1                      FALSE                      1
##   Latitude_MapAnything__c Longitude_MapAnything__c
## 1          26.79489          -80.06038
##   Library_Geolocation__Latitude__s Library_Geolocation__Longitude__s
## 1          43.45611          -95.31709
##   check_in_count
## 1              0
```

The MapAnything location:

The geolocation:

This library is actually displayed incorrectly in Iowa on the official webapp. The correct coordinates for this address are in Florida.

Another example is a location with a 101 degree difference in longitude.

```
libraries %>% filter(Latitude_MapAnything__c == 37.33889)
```

```
##      id      Name  Street__c
## 1 75881 LIB-000084125 North Campus Building, Perandori Dushan, Mitrovicë
##   City__c State_Province_Region__c Postal_Zip_Code__c  Country__c
## 1 Mitrovica                      KS          40000 United States
##   Traveling_Library__c Official_Charter_Number__c First_Map_Date__c Map_Me__c
## 1          FALSE                      85849      2022-10-13      TRUE
##   Map_Date__c Duplicate_Charter_Number__c Count_of_Primary_Stewards__c
## 1  2022-10-13                      FALSE                      1
##   Latitude_MapAnything__c Longitude_MapAnything__c
## 1          37.33889          -121.8825
##   Library_Geolocation__Latitude__s Library_Geolocation__Longitude__s
## 1          42.89542          20.86808
##   check_in_count
## 1              0
```




Figure 1: Lake Park in Florida




Figure 2: Lake Park in Iowa

The actual location is in Kosovo, but because they put “KS” as the state (which is Kansas, not Kosovo) this row was mistakenly assigned “United States” as its country.

The MapAnything location is in San Jose:




Figure 3: Location in San Jose

and the Geolocation is in Kosovo:

So we have two different examples where the correct coordinates are of different types. If we look at the distribution of coordinates we have:

```
libraries %>% select(Latitude_MapAnything__c, Library_Geolocation__Latitude__s, Longitude_MapAnything__c, Library_Geolocation__Longitude__s)

## Latitude_MapAnything__c Library_Geolocation__Latitude__s
## Min.      :-26.14      Min.      :-27.00
## 1st Qu.: 34.43      1st Qu.: 35.17
## Median : 39.58      Median : 39.74
## Mean   : 37.51      Mean   : 39.01
## 3rd Qu.: 42.44      3rd Qu.: 42.50
## Max.    : 71.30      Max.    : 86.95
## Longitude_MapAnything__c Library_Geolocation__Longitude__s
## Min.      :-170.47      Min.      :-170.49
## 1st Qu.: -105.06      1st Qu.: -105.96
## Median : -88.09      Median : -88.64
## Mean   : -89.96      Mean   : -93.56
```




Figure 4: Location in Kosovo

```
## 3rd Qu.: -79.91      3rd Qu.: -80.82
## Max.    : 77.56      Max.    : 158.07
```

Thus summary statistics are similar, but there are enough differences to cause concern. Note that there are a decent amount of rows where the MapAnything coordinates are (0, 0):

```
libraries %>% filter(Latitude_MapAnything__c == 0 & Longitude_MapAnything__c == 0) %>% count()

##          n
## 1 2428
```

One example is the library with id 14180.

```
libraries %>% filter(id == 14180)

##          id          Name      Street__c      City__c State_Province_Region__c
## 1 14180 LIB-000038331 11509 Kenny Dr Fort Worth TX
##      Postal_Zip_Code__c      Country__c      Traveling_Library__c
## 1          76244 United States          FALSE
##      Official_Charter_Number__c First_Map_Date__c Map_Me__c Map_Date__c
## 1          77029      2019-03-21      TRUE      2023-04-07
##      Duplicate_Charter_Number__c Count_of_Primary_Stewards__c
## 1          FALSE          1
##      Latitude_MapAnything__c Longitude_MapAnything__c
## 1          0          0
##      Library_Geolocation__Latitude__s Library_Geolocation__Longitude__s
## 1          32.93978          -97.27757
##      check_in_count
## 1          1
```

None of the values look notable other than the (0, 0) MapAnything coordinates, and this library shows up on the official map.

These (0, 0) coordinates are basically missing values as all the libraries we are looking at are located in the US so (0, 0) is definitely an invalid coordinate. If we look at the webapp, it appears that the developers use the geolocation values on the interactive map:

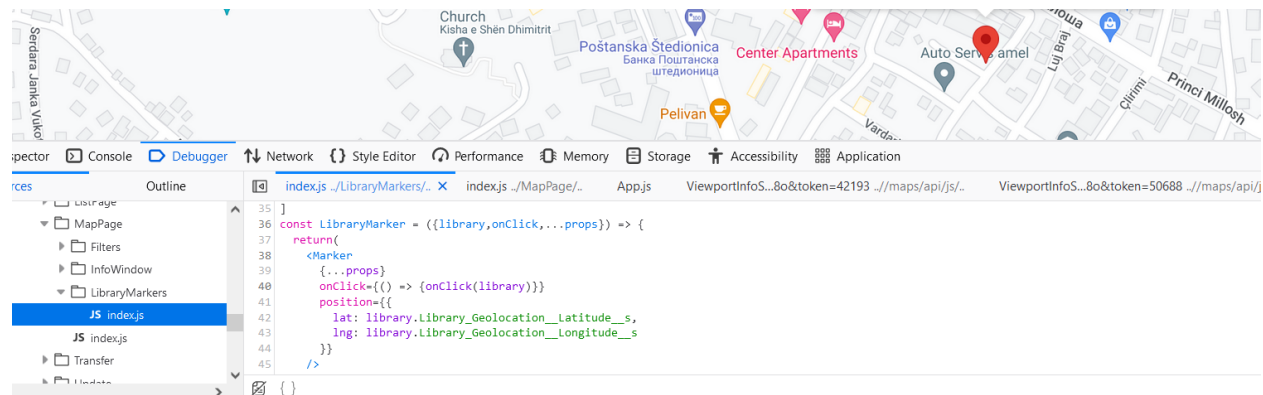


Figure 5: Screenshot of javascript snippet using Library_Geolocation for the pins on the map

So let's default to the geolocation coordinates, but use the MapAnything coordinates if they match up with the address listed. Let's then drop every data point with coordinates not located within the United States. In order to determine this we can look at what U.S. state a given library is in and then filter out the libraries with no state values.

```

# Convert the coordinates to a sf object
# Our coordinate reference system is the WGS84 standard which is what Google
# maps uses. Its EPSG Code is 4326. The format for a point is (longitude, latitude).
lib_pts <- libraries %>% st_as_sf(coords = c("Library_Geolocation__Longitude__s", "Library_Geolocation__Latitude__s"))

# For comparison also convert the MapAnything coordinates into a sf object
lib_pts_alt <- libraries %>% st_as_sf(coords = c("Longitude_MapAnything__c", "Latitude_MapAnything__c"))

# Read in and transform the GADM data to WGS84 format.
GADM_data <- st_read(dsn = "gadm36_USA_gpkg/gadm36_USA.gpkg", layer = "gadm36_USA_1")

## Reading layer `gadm36_USA_1' from data source
##   `/Users/kcrans/Desktop/projects/little_free/gadm36_USA_gpkg/gadm36_USA.gpkg'
##   using driver `GPKG'
## Simple feature collection with 51 features and 10 fields
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -179.1506 ymin: 18.90986 xmax: 179.7734 ymax: 72.6875
## Geodetic CRS:   WGS 84

state_pts <- st_transform(GADM_data, crs = 4326)

# Make a data.frame of all the possible state names
state_names <- state_pts$NAME_1

# Find the intersections between the library points and state polygons
# Convert to an integer to use as an index in the state names data.frame.
classifications <- as.integer(st_intersects(lib_pts, state_pts))
alt_classifications <- as.integer(st_intersects(lib_pts_alt, state_pts))

# Store the results in a copy of libraries
temp_libraries <- libraries %>% mutate(
  state = state_names[classifications], alt_state = state_names[alt_classifications])
sum(is.na(temp_libraries$state) & is.na(temp_libraries$alt_state))

## [1] 70

```

There are 70 locations with (both) coordinates not in the U.S. for whatever reason. Let's take a look at a few of them.

```
temp_libraries %>% filter(is.na(state) & is.na(alt_state)) %>% head()
```

##	id	Name	Street__c	City__c
## 1	2278	LIB-000026204	Perumkulam	Kottarakara
## 2	3919	LIB-000027852	Keerkring 112	De Meern
## 3	4141	LIB-000028075	Aljazar st	Khartoum
## 4	4515	LIB-000028452	Chequers Corner, Hurst Drive	Walton on the Hill
## 5	5604	LIB-000028915	1 Stirrup Close	Newbury
## 6	7355	LIB-000030684	Dorfstrasse 21	Bühler
##	State_Province_Region__c	Postal_Zip_Code__c	Country__c	
## 1	<NA>	691566	United States	
## 2	<NA>	3454kz	United States	
## 3	<NA>	11112	United States	
## 4	<NA>	KT20 7QT	United States	
## 5	<NA>	RG14 7XD	United States	
## 6	<NA>	9055	United States	

```

##   Traveling_Library__c Official_Charter_Number__c First_Map_Date__c Map_Me__c
## 1                FALSE                57110        2017-06-28        TRUE
## 2                FALSE                58388        2017-08-29        TRUE
## 3                FALSE                46879        2017-09-07        TRUE
## 4                FALSE                46255        2017-09-28        TRUE
## 5                FALSE                60235        2017-10-16        TRUE
## 6                FALSE                M68171        2018-01-29        TRUE
##   Map_Date__c Duplicate_Charter_Number__c Count_of_Primary_Stewards__c
## 1 2017-06-28                FALSE                1
## 2 2017-08-29                FALSE                1
## 3 2017-09-07                FALSE                1
## 4 2017-09-28                FALSE                1
## 5 2017-10-16                FALSE                1
## 6 2018-02-19                FALSE                2
##   Latitude_MapAnything__c Longitude_MapAnything__c
## 1          9.03882          76.7597700
## 2         52.08682           5.0297300
## 3         15.50065          32.5598994
## 4         51.27481          -0.2506241
## 5         51.38518          -1.3159300
## 6         47.37292           9.4252400
##   Library_Geolocation__Latitude__s Library_Geolocation__Longitude__s
## 1          9.03882          76.7597700
## 2         52.08682           5.0297300
## 3         15.50065          32.5598994
## 4         51.27481          -0.2506241
## 5         51.38518          -1.3159300
## 6         47.37292           9.4252400
##   check_in_count state alt_state
## 1             0 <NA>    <NA>
## 2             0 <NA>    <NA>
## 3             0 <NA>    <NA>
## 4             0 <NA>    <NA>
## 5             0 <NA>    <NA>
## 6             0 <NA>    <NA>

```

Some of the rows are for foreign libraries, but it look like the majority are libraries with no street entries. Some like charter number G10014(148 Marina Plaza Dunedin) are located very close to the ocean and hence were classified incorrectly due to the resolution of the geography in the GADM dataset. A few like 150219 (1710 S Trenon Ave Tulsa) are mislabeled with coordinates not in the Unites States. By manual inspection, it looks the MapAnything coordinates give the appropriate state for some of the cases where the locations are right near a body of water or the street address is missing. When a library is located in a country other than the United States, both columns will have NA values. In general, if we take a look at the state assignments based off the geolocation coordinates (“state”) and MapAnything coordinates (“alt_state”), the possibilities are:

1. NA for both. This means either the library is not in the U.S and should be removed from our dataframe, or it is located in the U.S. but too close to the ocean. We filter out data in the former instance and use the “State_Province_Region__c” assignment in the latter.
2. An actual state for the geolocation coordinates and NA for MapAnything. We should use the “state” assignment then.
3. An actual state for the MapAnything coordinates and NA for the geolocation. We will go with the “alt_state” assignment in this case.

4. and 5. Actual states for both coordinates. If they are the same, we will go with that assignment. If they are different, we will use whichever assignment lines up with the “State_Province_Region__c” value for that datapoint.

Now, let’s create a new column `state_name` to hold whatever state name based off the above criteria we choose, and also `long` and `lat` columns to hold the coordinates we end up using.

First, here’s a helper function to convert state abbreviations into full names:

```
convert_state <- function(state_code) {  
  up_code <- toupper(state_code) # Ignore case differences  
  name_str <- switch(up_code,  
    "AL" = "Alabama",  
    "AK" = "Alaska",  
    "AZ" = "Alaska",  
    "AR" = "Arkansas",  
    "AS" = "American Samoa",  
    "CA" = "California",  
    "CO" = "Colorado",  
    "CT" = "Connecticut",  
    "DE" = "Delaware",  
    "DC" = "District of Columbia",  
    "FL" = "Florida",  
    "GA" = "Georgia",  
    "GU" = "Guam",  
    "HI" = "Hawaii",  
    "ID" = "Idaho",  
    "IL" = "Illinois",  
    "IN" = "Indiana",  
    "IA" = "Iowa",  
    "KS" = "Kansas",  
    "KY" = "Kentucky",  
    "LA" = "Louisiana",  
    "ME" = "Maine",  
    "MD" = "Maryland",  
    "MA" = "Massachusetts",  
    "MI" = "Michigan",  
    "MN" = "Minnesota",  
    "MS" = "Mississippi",  
    "MO" = "Missouri",  
    "MT" = "Montana",  
    "NE" = "Nebraska",  
    "NV" = "Nevada",  
    "NH" = "New Hampshire",  
    "NJ" = "New Jersey",  
    "NM" = "New Mexico",  
    "NY" = "New York",  
    "NC" = "North Carolina",  
    "ND" = "North Dakota",  
    "MP" = "Northern Mariana Islands",  
    "OH" = "Ohio",  
    "OK" = "Oklahoma",  
    "OR" = "Oregon",  
    "PA" = "Pennsylvania",  
    "PR" = "Puerto Rico",  
  )  
}
```



```

"RI" = "Rhode Island",
"SC" = "South Carolina",
"SD" = "South Dakota",
"TN" = "Tennessee",
"TX" = "Texas",
"TT" = "Trust Territories",
"UT" = "Utah",
"VT" = "Vermont",
"VA" = "Virginia",
"VI" = "Virgin Islands",
"WA" = "Washington",
"WV" = "West Virginia",
"WI" = "Wisconsin",
"WY" = "Wyoming",
state_code # Else return the string as-is
)
return(name_str)
}
convert_state <- Vectorize(convert_state)

```

Let's build our dataframe case-by-case:

1.

```

temp_libraries %>% filter(is.na(state) & is.na(alt_state)) %>% filter(!is.na(State_Province_Region__c))

```

##	id	Name	Street__c	City__c
## 1	10894 LIB-000034171	1932 S. Oceanshore Blvd		Flagler Beach
## 2	19119 LIB-000046571	319 Dunham Point Road		Deer Isle
## 3	19356 LIB-000043089		<NA>	Alytus
## 4	26622 LIB-000053055		1 Newport Ave	North Kingstown
## 5	34172 LIB-000064463		567 Angell Street	Providence
## 6	46526 LIB-000009684		510 Shorewood Drive	International Falls
## 7	51495 LIB-000014854		86 Butts Rock Rd.	Little Compton
## 8	51556 LIB-000014915		Ocean Ave & Lincoln	Avon by the Sea
## 9	56156 LIB-000019673		2013 Wildwood Lane	Anchorage
## 10	58958 LIB-000022491		1800 gulf road	Tarpon springs
## 11	63301 LIB-000071612		185 Ferry Road	Saunderstown
## 12	64910 LIB-000073170		261 Hart Street	Dighton
## 13	75691 LIB-000083936		22 Beachwood Drive	Warwick
## 14	78630 LIB-000086854		<NA>	Sandbridge
## 15	78786 LIB-000087009		<NA>	Madison
## 16	79239 LIB-000087459	114 W. Chicago Street		Caldwell
## 17	82395 LIB-000090612	64 Sleeper Street		Boston
##	State_Province_Region__c	Postal_Zip_Code__c	Country__c	
## 1	FL	32136	United States	
## 2	ME	04627	United States	
## 3	AL	63210	United States	
## 4	CT	02852	United States	
## 5	RI	02906	United States	
## 6	MN	56649	United States	
## 7	RI	02837	United States	
## 8	NJ	07717	United States	
## 9	AK	99517	United States	
## 10	FL	34689	United States	

## 11	RI	02874 United States		
## 12	MA	02715 United States		
## 13	RI	02818 United States		
## 14	VA	23456 United States		
## 15	MS	39110 United States		
## 16	ID	83605 United States		
## 17	MA	02210 United States		
##	Traveling_Library__c	Official_Charter_Number__c	First_Map_Date__c	Map_Me__c
## 1	FALSE	76585	2018-08-03	TRUE
## 2	FALSE	83568	2020-04-13	TRUE
## 3	FALSE	13135	2019-10-01	TRUE
## 4	FALSE	107184	2020-09-02	TRUE
## 5	FALSE	117921	2021-05-03	TRUE
## 6	FALSE	6426	2014-10-10	TRUE
## 7	FALSE	26399	2015-09-10	TRUE
## 8	FALSE	29517	2015-09-10	TRUE
## 9	FALSE	39053	2016-07-05	TRUE
## 10	FALSE	23484	2016-11-14	TRUE
## 11	FALSE	135342	2021-10-13	TRUE
## 12	FALSE	135518	2021-12-13	TRUE
## 13	FALSE	152156	2022-10-04	TRUE
## 14	FALSE	131731	2023-02-21	TRUE
## 15	FALSE	157744	2023-03-01	TRUE
## 16	FALSE	162828	2023-03-20	TRUE
## 17	FALSE	160393	2023-04-10	TRUE
##	Map_Date__c	Duplicate_Charter_Number__c	Count_of_Primary_Stewards__c	
## 1	2018-08-03	FALSE	1	
## 2	2020-04-13	FALSE	1	
## 3	2019-10-01	FALSE	1	
## 4	2020-09-02	FALSE	1	
## 5	2021-05-03	FALSE	1	
## 6	2019-01-22	FALSE	1	
## 7	2020-04-01	FALSE	1	
## 8	2015-09-10	FALSE	1	
## 9	2016-07-05	FALSE	1	
## 10	2016-11-14	FALSE	1	
## 11	2021-10-18	FALSE	1	
## 12	2021-12-13	FALSE	1	
## 13	2022-10-04	FALSE	1	
## 14	2023-02-21	FALSE	1	
## 15	2023-03-01	FALSE	1	
## 16	2023-03-20	FALSE	1	
## 17	2023-04-10	FALSE	1	
##	Latitude_MapAnything__c	Longitude_MapAnything__c		
## 1	29.46048	-81.11792		
## 2	44.21357	-68.71957		
## 3	54.40273	24.03162		
## 4	41.57126	-71.44160		
## 5	41.83050	-71.38526		
## 6	48.58890	-93.46442		
## 7	41.48505	-71.14141		
## 8	40.18972	-74.00997		
## 9	61.20429	-149.92247		
## 10	28.14467	-82.78962		


```
## 11          41.50778          -71.41860
## 12          41.80444          -71.12383
## 13          41.66889          -71.42342
## 14          0.00000          0.00000
## 15          0.00000          0.00000
## 16          0.00000          0.00000
## 17          0.00000          0.00000
##      Library_Geolocation__Latitude__s Library_Geolocation__Longitude__s
## 1          29.46050          -81.11795
## 2          44.21276          -68.71880
## 3          54.40512          24.02948
## 4          41.57137          -71.44158
## 5          41.83050          -71.38526
## 6          48.58890          -93.46443
## 7          41.48536          -71.14151
## 8          40.18918          -74.01003
## 9          61.20386          -149.92313
## 10         28.14444          -82.78960
## 11         41.50778          -71.41860
## 12         41.80444          -71.12383
## 13         41.66889          -71.42342
## 14         36.42580          -75.56300
## 15         32.46640          90.15027
## 16         43.67455          116.69132
## 17         42.35282          -71.04923
##      check_in_count state alt_state
## 1          2 <NA> <NA>
## 2          0 <NA> <NA>
## 3          0 <NA> <NA>
## 4          4 <NA> <NA>
## 5         17 <NA> <NA>
## 6          0 <NA> <NA>
## 7          1 <NA> <NA>
## 8          0 <NA> <NA>
## 9          0 <NA> <NA>
## 10         1 <NA> <NA>
## 11         4 <NA> <NA>
## 12         2 <NA> <NA>
## 13         2 <NA> <NA>
## 14         0 <NA> <NA>
## 15         0 <NA> <NA>
## 16         0 <NA> <NA>
## 17         0 <NA> <NA>
```

The library in Alytus, Lithuania is the only one not in the U.S. So let's just drop that one and add the new state assignments:

```
case_1 <- temp_libraries %>% filter(is.na(state) & is.na(alt_state)) %>% filter(!is.na(State_Province_Region__c)) %>%
  filter(City__c != "Alytus") %>% mutate(state_name = convert_state(State_Province_Region__c), long = Library_Geolocation__Longitude__s, lat = Library_Geolocation__Latitude__s)
```

2.

```
case_2 <- temp_libraries %>% filter(!is.na(state) & is.na(alt_state)) %>%
  mutate(state_name = state, long = Library_Geolocation__Longitude__s, lat = Library_Geolocation__Latitude__s)
```

3.

```
case_3 <- temp_libraries %>% filter(is.na(state) & !is.na(alt_state)) %>%
  mutate(state_name = alt_state, long = Longitude_MapAnything_c, lat = Latitude_MapAnything_c)
```

4.

```
case_4 <- temp_libraries %>% filter(state == alt_state) %>% mutate(state_name = state, long = Longitude_MapAnything_c, lat = Latitude_MapAnything_c)
```

5.

```
# state != alt_state but the given state code is equal to state
case_5a <- temp_libraries %>% filter(state != alt_state ) %>% mutate(State_Province_Region__c = convert_to_state_province_region(state))
```

```
# state != alt_state but the given state code is equal to alt_state
case_5b <- temp_libraries %>% filter(state != alt_state ) %>% mutate(State_Province_Region__c = convert_to_state_province_region(alt_state))
```


Joining it all together:

```
temp_libraries <- bind_rows(case_1, case_2, case_3, case_4, case_5a, case_5b)
# Drop state, alt_state, and both original coordinates
libraries <- select(temp_libraries, -c(state, alt_state, State_Province_Region__c, Library_Geolocation__c))
arrange(id) # Sort rows based off id
```


Finally we can do some analysis with states:

```
state_counts <- libraries %>% group_by(state_name) %>% count %>% arrange(desc(n))
state_counts <- state_counts %>% rename("NAME" = "state_name")
```

```
state_counts %>% ggplot(aes(x = n)) +
  geom_histogram(binwidth = 1000)
```



```
st_transform(us_states, crs = 3857) %>% full_join(state_counts, by = join_by(NAME)) %>%
  ggplot(aes(fill = n)) +
  geom_sf() +
  labs(fill = "Number of LFLs") +
  scale_fill_fermenter(n.breaks = 9, direction = 1, palette = "Oranges") +
  ggtitle("Number of Little Free Libraries per state")
```



```
#ggsave("lflState.png")
```

And then if we normalize by population:

```
state_pops = as_tibble(read_csv("census_pop_data.csv"))
```


```
## Rows: 71 Columns: 53
## -- Column specification -----
## Delimiter: ","
## chr (1): Label (Grouping)
## num (52): Alabama, Alaska, Arizona, Arkansas, California, Colorado, Connecti...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
state_pops <- state_pops[1,] %>% select(-"Label (Grouping)" ) %>% pivot_longer(cols = everything(), nam
norm_state_counts <- state_counts %>% left_join(state_pops) %>% mutate(per_cap = 10000*n/pop) %>% select
```

```
## Joining with `by = join_by(NAME)`
```

```
st_transform(us_states, crs = 3857) %>% left_join(norm_state_counts, by = join_by(NAME)) %>%
  ggplot(aes(fill = per_cap)) +
  geom_sf() +
  labs(fill = "LFLs per 10,000 residents") +
  scale_fill_fermenter(n.breaks = 9, direction = 1, palette = "Oranges") +
  ggtitle("Number of Little Free Libraries (population normalized)")
```

Number of Little Free Libraries (population normalized)



```
#ggsave("lflStateNorm.png")
```

Per capita, it looks like the distribution of LFLs is biased towards the Midwest, in particular Minnesota and Wisconsin. This is not surprising, as the organization was founded in Hudson, Wisconsin and eventually moved to Minneapolis, Minnesota.

Analysis by county

There are many faults to doing geospatial analysis based off zipcodes (they overlap, are not contiguous, etc...). Moreover, in the most recent 2020 census differential privacy techniques applied to data on finer scales resulted in a lot of inaccuracies. So for the purposes of this analysis let's use county (and county-equivalent) data.

```
library(tidycensus)
census_api_key("7ce885ce139c2116d8d26ffca665df473f98a98c")

## To install your API key for use in future sessions, run this function with `install = TRUE`.
county_incomes <- get_acs(geography = "county", variables = "B19013_001", year = 2021)

## Getting data from the 2017-2021 5-year ACS
county_populations <- get_acs(geography = "county", variables = "B02001_001", year = 2021)

## Getting data from the 2017-2021 5-year ACS
county_incomes <- county_incomes %>% full_join(county_populations, by = "NAME")

# Convert the coordinates to a sf object
# Our coordinate reference system is the WGS84 standard which is what Google
# maps uses. Its EPSG Code is 4326. The format for a point is (longitude, latitude).
lib_pts <- libraries %>% st_as_sf(coords = c("long", "lat"), crs = st_crs(4326))

# Read in and transform the GADM county data to WGS84 format.
GADM_data2 <- st_read(dsn = "gadm36_USA_gpkg/gadm36_USA.gpkg", layer = "gadm36_USA_2")

## Reading layer `gadm36_USA_2` from data source
##   `/Users/kcrans/Desktop/projects/little_free/gadm36_USA_gpkg/gadm36_USA.gpkg'
##   using driver `GPKG'
## Simple feature collection with 3148 features and 13 fields
## Geometry type: MULTIPOLYGON
```

```
## Dimension:      XY
## Bounding box:   xmin: -179.1506 ymin: 18.90986 xmax: 179.7734 ymax: 72.6875
## Geodetic CRS:   WGS 84

county_pts <- st_transform(GADM_data2, crs = 4326)

# Differences between GADM and 2021 ACS county definitions means we have to do a lot of manual data cleaning

county_pts <- county_pts %>% mutate(COUNTY = recode(TYPE_2, "Independent City" = "city", "City And County" = "city"),
  mutate(NAME = recode(NAME_2, "Fairfax City" = "Fairfax", "Carson City" = "Carson", "Roanoke City" = "Roanoke"),
  mutate(NAME = str_replace_all(NAME, "Saint", "St.)) %>%
  mutate(NAME = str_replace_all(NAME, "De Kalb", "DeKalb")) %>%
  mutate(NAME = str_replace_all(NAME, "La Salle", "LaSalle")) %>%
  unite(county_name, c("NAME", "COUNTY"), sep = " ") %>% mutate(county_name = recode(county_name, "Cliff County" = "Cliff"))

# Make a data.frame of all the possible county names
county_names <- county_pts$county_name

# Find the intersections between the library points and state polygons
# Convert to an integer to use as an index in the state names data.frame.
classifications <- as.integer(st_intersects(lib_pts, county_pts))

temp_libraries <- libraries %>% mutate(
  county = county_names[classifications])
```

These are all the libraries which did not intersect any of the county boundaries we looked at.

```
temp_libraries %>% filter(is.na(county))
```

##	id	Name	Street__c	City__c
## 1	10894	LIB-000034171	1932 S. Oceanshore Blvd	Flagler Beach
## 2	19119	LIB-000046571	319 Dunham Point Road	Deer Isle
## 3	26622	LIB-000053055	1 Newport Ave	North Kingstown
## 4	34172	LIB-000064463	567 Angell Street	Providence
## 5	46526	LIB-000009684	510 Shorewood Drive	International Falls
## 6	51495	LIB-000014854	86 Butts Rock Rd.	Little Compton
## 7	51556	LIB-000014915	Ocean Ave & Lincoln	Avon by the Sea
## 8	56156	LIB-000019673	2013 Wildwood Lane	Anchorage
## 9	58958	LIB-000022491	1800 gulf road	Tarpon springs
## 10	63301	LIB-000071612	185 Ferry Road	Saunderstown
## 11	64910	LIB-000073170	261 Hart Street	Dighton
## 12	75691	LIB-000083936	22 Beachwood Drive	Warwick
## 13	78630	LIB-000086854	<NA>	Sandbridge
## 14	78786	LIB-000087009	<NA>	Madison
## 15	79239	LIB-000087459	114 W. Chicago Street	Caldwell
## 16	82395	LIB-000090612	64 Sleeper Street	Boston
##	Postal_Zip_Code__c	Country__c	Traveling_Library__c	
## 1	32136	United States	FALSE	
## 2	04627	United States	FALSE	
## 3	02852	United States	FALSE	
## 4	02906	United States	FALSE	
## 5	56649	United States	FALSE	
## 6	02837	United States	FALSE	
## 7	07717	United States	FALSE	
## 8	99517	United States	FALSE	

## 9	34689	United States	FALSE	
## 10	02874	United States	FALSE	
## 11	02715	United States	FALSE	
## 12	02818	United States	FALSE	
## 13	23456	United States	FALSE	
## 14	39110	United States	FALSE	
## 15	83605	United States	FALSE	
## 16	02210	United States	FALSE	
##	Official_Charter_Number__c	First_Map_Date__c	Map_Me__c	Map_Date__c
## 1	76585	2018-08-03	TRUE	2018-08-03
## 2	83568	2020-04-13	TRUE	2020-04-13
## 3	107184	2020-09-02	TRUE	2020-09-02
## 4	117921	2021-05-03	TRUE	2021-05-03
## 5	6426	2014-10-10	TRUE	2019-01-22
## 6	26399	2015-09-10	TRUE	2020-04-01
## 7	29517	2015-09-10	TRUE	2015-09-10
## 8	39053	2016-07-05	TRUE	2016-07-05
## 9	23484	2016-11-14	TRUE	2016-11-14
## 10	135342	2021-10-13	TRUE	2021-10-18
## 11	135518	2021-12-13	TRUE	2021-12-13
## 12	152156	2022-10-04	TRUE	2022-10-04
## 13	131731	2023-02-21	TRUE	2023-02-21
## 14	157744	2023-03-01	TRUE	2023-03-01
## 15	162828	2023-03-20	TRUE	2023-03-20
## 16	160393	2023-04-10	TRUE	2023-04-10
##	Duplicate_Charter_Number__c	Count_of_Primary_Stewards__c	check_in_count	
## 1	FALSE	1	2	
## 2	FALSE	1	0	
## 3	FALSE	1	4	
## 4	FALSE	1	17	
## 5	FALSE	1	0	
## 6	FALSE	1	1	
## 7	FALSE	1	0	
## 8	FALSE	1	0	
## 9	FALSE	1	1	
## 10	FALSE	1	4	
## 11	FALSE	1	2	
## 12	FALSE	1	2	
## 13	FALSE	1	0	
## 14	FALSE	1	0	
## 15	FALSE	1	0	
## 16	FALSE	1	0	
##	state_name	long	lat	county
## 1	Florida	-81.11795	29.46050	<NA>
## 2	Maine	-68.71880	44.21276	<NA>
## 3	Connecticut	-71.44158	41.57137	<NA>
## 4	Rhode Island	-71.38526	41.83050	<NA>
## 5	Minnesota	-93.46443	48.58890	<NA>
## 6	Rhode Island	-71.14151	41.48536	<NA>
## 7	New Jersey	-74.01003	40.18918	<NA>
## 8	Alaska	-149.92313	61.20386	<NA>
## 9	Florida	-82.78960	28.14444	<NA>
## 10	Rhode Island	-71.41860	41.50778	<NA>
## 11	Massachusetts	-71.12383	41.80444	<NA>

```
## 12 Rhode Island -71.42342 41.66889 <NA>
## 13 Virginia -75.56300 36.42580 <NA>
## 14 Mississippi 90.15027 32.46640 <NA>
## 15 Idaho 116.69132 43.67455 <NA>
## 16 Massachusetts -71.04923 42.35282 <NA>
```

These all appear to be locations too close to maritime borders for the resolution in our geometry objects.

Manually fix all the county classifications that have changed (i.e. new counties and consolidations)

```
libraries <- temp_libraries %>% unite(full_name, c("county", "state_name"), sep = ", ") %>%
  mutate(full_name = recode(full_name, "Shannon County, South Dakota" = "Oglala Lakota County, South Dakota")) %>%
  mutate(full_name = replace(full_name, City__c == "Valdez", "Chugach Census Area, Alaska")) %>%
  mutate(full_name = replace(full_name, City__c == "Kake", "Prince of Wales-Hyder Census Area, Alaska")) %>%
  mutate(full_name = replace(full_name, City__c == "Wrangell", "Wrangell City and Borough, Alaska")) %>%
  mutate(full_name = replace(full_name, City__c == "Sackets Harbor", "Jefferson County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Washburn" & Postal_Zip_Code__c == "54891", "Bayfield County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Baileys Harbor", "Door County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Duluth" & Postal_Zip_Code__c == "55802", "St. Louis County, Minnesota")) %>%
  mutate(full_name = replace(full_name, City__c == "Sodus Pt.", "Wayne County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Bridgman", "Berrien County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Menominee", "Menominee County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Erie" & Postal_Zip_Code__c == "16505", "Erie County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Niagara Falls", "Niagara County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Sault Ste. Marie", "Chippewa County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Sturgeon Bay", "Door County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Manistee", "Manistee County, Michigan")) %>%
  mutate(full_name = replace(full_name, Postal_Zip_Code__c == "14612", "Monroe County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Houghton", "Houghton County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Mears", "Oceana County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Racine", "Racine County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Lake Linden", "Houghton County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Muskegon", "Muskegon County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "East Tawas", "Iosco County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Fineview", "Jefferson County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Augres", "Arenac County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Petoskey", "Emmet County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Omena", "Leelanau County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Sheboygan", "Sheboygan County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Houghton", "Houghton County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "MANITOWOC", "Manitowoc County, Wisconsin")) %>%
  mutate(full_name = replace(full_name, City__c == "Port Huron", "St. Clair County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Detroit", "Wayne County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Wolcott" & Postal_Zip_Code__c == "14590", "Wayne County, New York")) %>%
  mutate(full_name = replace(full_name, City__c == "Tawas City", "Iosco County, Michigan")) %>%
  mutate(full_name = replace(full_name, City__c == "Marinette", "Marinette County, Wisconsin")) %>%

  mutate(full_name = replace(full_name, City__c == "Flagler Beach", "Flagler County, Florida")) %>%
  mutate(full_name = replace(full_name, City__c == "Deer Isle", "Hancock County, Maine")) %>%
  mutate(full_name = replace(full_name, City__c == "North Kingstown", "Washington County, Rhode Island")) %>%
  mutate(full_name = replace(full_name, City__c == "North Kingstown", "Washington County, Rhode Island")) %>%
  mutate(full_name = replace(full_name, City__c == "Providence" & Postal_Zip_Code__c == "02906", "Providence County, Rhode Island")) %>%
  mutate(full_name = replace(full_name, City__c == "International Falls", "Koochiching County, Minnesota")) %>%
  mutate(full_name = replace(full_name, City__c == "Little Compton", "Newport County, Rhode Island")) %>%
  mutate(full_name = replace(full_name, City__c == "Avon by the Sea", "Monmouth County, New Jersey")) %>%
```

```

mutate(full_name = replace(full_name, City__c == "Anchorage", "Jefferson County, Kentucky")) %>%
mutate(full_name = replace(full_name, City__c == "Tarpon springs", "Pinellas County, Florida")) %>%
mutate(full_name = replace(full_name, City__c == "Saunderstown", "Washington County, Rhode Island")) %>%
mutate(full_name = replace(full_name, City__c == "Dighton", "Bristol County, Massachusetts")) %>%
mutate(full_name = replace(full_name, City__c == "Warwick" & Postal_Zip_Code__c == "02818", "Kent County, Rhode Island")) %>%
mutate(full_name = replace(full_name, City__c == "Sandbridge", "Virginia Beach city, Virginia")) %>%
mutate(full_name = replace(full_name, City__c == "Madison" & Postal_Zip_Code__c == "39110", "Madison County, Georgia")) %>%
mutate(full_name = replace(full_name, City__c == "Caldwell" & Postal_Zip_Code__c == "83605", "Canyon County, Idaho")) %>%
mutate(full_name = replace(full_name, City__c == "Boston" & Postal_Zip_Code__c == "02210", "Suffolk County, Massachusetts")) %>%
mutate(NAME = tolower(full_name))

```

```

county_counts <- libraries %>% group_by(NAME) %>% count %>% arrange(desc(n))
county_incomes <- county_incomes %>% mutate(NAME = tolower(NAME))
county_stats <- county_counts %>% left_join(county_incomes, by = "NAME")

```

```

county_stats <- rename(county_stats, est_income = estimate.x, est_pop = estimate.y, moe = moe.x ) %>% select(-moe)

```

Top 5 counties by total number of little free libraries:

```

county_stats %>% arrange(desc(n)) %>% head(5)

```

```

## # A tibble: 5 x 5
## # Groups:   NAME [5]
##   NAME                                n est_income est_pop libs_percap
##   <chr>                                <int>     <dbl>    <dbl>      <dbl>
## 1 los angeles county, california    1337      76367  10019635      1.33
## 2 king county, washington           1062     106326   2240876      4.74
## 3 cook county, illinois              918      72121   5265398      1.74
## 4 hennepin county, minnesota         912      85438   1270283      7.18
## 5 san diego county, california       799      88240   3296317      2.42

```

Top 5 counties by number of little free libraries adjusted for population:

```

county_stats %>% arrange(desc(libs_percap)) %>% head(5)

```

```

## # A tibble: 5 x 5
## # Groups:   NAME [5]
##   NAME                                n est_income est_pop libs_percap
##   <chr>                                <int>     <dbl>    <dbl>      <dbl>
## 1 sherman county, oregon              4      53606    1784      22.4
## 2 daggett county, utah                 1      49792     564      17.7
## 3 lewis county, idaho                   6      44028    3613      16.6
## 4 jerauld county, south dakota          3      58375    1811      16.6
## 5 banner county, nebraska              1      57917     605      16.5

```

```

summary(lm(libs_percap ~ est_income, data = county_stats))

```


```

##
## Call:
## lm(formula = libs_percap ~ est_income, data = county_stats)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7334 -1.1971 -0.5822  0.5199 20.3427
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)

```



```
## (Intercept) 1.165e+00 1.588e-01 7.337 2.97e-13 ***
## est_income 1.705e-05 2.544e-06 6.701 2.57e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.994 on 2438 degrees of freedom
## Multiple R-squared:  0.01808,    Adjusted R-squared:  0.01768
## F-statistic: 44.9 on 1 and 2438 DF,  p-value: 2.569e-11
plot(libs_percap ~ est_income, data = county_stats)
abline(lm(libs_percap ~ est_income, data = county_stats))
```




So there is a low correlation between average income and number of little free libraries per capita. If we ignore the adjustment for population, we get:

```
summary(lm(n ~ est_income, data = county_stats))

##
## Call:
## lm(formula = n ~ est_income, data = county_stats)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -157.53  -20.55   -8.02    4.37  1288.32
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.567e+01  5.208e+00  -12.61  <2e-16 ***
## est_income   1.497e-03  8.342e-05   17.95  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 65.37 on 2438 degrees of freedom
## Multiple R-squared:  0.1167, Adjusted R-squared:  0.1164
```

```
## F-statistic: 322.2 on 1 and 2438 DF,  p-value: < 2.2e-16
```

```
plot(n ~ est_income, data = county_stats)  
abline(lm(n ~ est_income, data = county_stats))
```



Note that incomes can vary significantly within a county (there can be rich and poor cities, rich and poor neighborhoods, etc...). The analysis would be significantly improved if we used income figures on a finer-grained scale.